

# Query/Answer Strategies Comparison

"The search for the perfect language is the story of a dream and of a series of failures."  
— Umberto Eco, *The Search for the Perfect Language*

**Historical note:** Leibniz's *characteristica universalis* assigned prime numbers to primitive concepts — "animal" = 2, "rational" = 3 — so that composite concepts became products: "human" =  $2 \times 3 = 6$ . Reasoning reduced to divisibility tests. Modern embeddings pursue the same dream: mapping meaning to vectors where semantic relationships become geometric operations (cosine similarity). The difference is that Leibniz required *a priori* enumeration of all primitive ideas, while embeddings learn structure statistically from data.


## OpenRouter API Token:

```
k-or-v1-c2b3fed13934605c6b75e9962fbef4b867bf2da694d1661f5d88229cda0da3c
```

## Acronyms & Terms:

- **RAG** — Retrieval-Augmented Generation: combining retrieval with LLM generation
- **BM25** — Best Matching 25: classic keyword-based ranking algorithm (TF-IDF variant)
- **RAPTOR** — Recursive Abstractive Processing for Tree-Organized Retrieval
- **CoIPali** — Vision-language model for visual document embeddings (based on PaliGemma)
- **LLM** — Large Language Model
- **VLM** — Vision-Language Model
- **Cross-encoder** — Reranker model that scores query-document pairs jointly


## Strategy 1: Small Router Model



Routes query to most relevant document(s) using a lightweight classifier, then reranks before LLM.

# Understanding Word Vectors

Word embeddings map words to vectors in high-dimensional space, where **semantic relationships become geometric relationships**. The classic example: the vector from "man" to "woman" captures the concept of gender — and this same vector connects "king" to "queen".




**The famous equation:**  $\text{king} - \text{man} + \text{woman} \approx \text{queen}$

This works because embeddings encode *relational structure*. The "royalty" dimension is preserved while the "gender" dimension shifts. This is why vector similarity (cosine distance) can capture meaning — semantically related words cluster together, and analogies become vector arithmetic.

**From words to documents:** RAG systems extend this idea. Instead of single words, we embed chunks or entire pages, then find the most similar vectors to a query. The same geometric intuition applies — relevant documents are "nearby" in embedding space.

## Strategy 2: RAG + BM25 + Visual Embeddings

**Strategy 2: RAG + BM25 + Visual Embeddings**




*Three parallel paths: text embeddings (semantic), visual embeddings (ColPali — page as image), and BM25 (keyword). Results merged and reranked.*

# Strategy 3: RAPTOR

"The defect of all symbolic systems lies in their inability to account for context." — Umberto Eco

## Strategy 3: RAPTOR (Hierarchical Summarization)



*Builds a hierarchical tree of summaries. Query searches all levels — chunks, summaries, and abstract root.*

# Strengths & Weaknesses Comparison

Strategy	Strengths	Weaknesses	Best For
1. Router (Small Model)	<ul style="list-style-type: none"> <li>• Very fast latency</li> <li>• Low compute cost</li> <li>• Simple to implement</li> <li>• Easy to debug</li> </ul>	<ul style="list-style-type: none"> <li>• Loses within-doc granularity</li> <li>• May need retraining</li> <li>• Binary doc selection</li> </ul>	Clearly separable document topics
2. RAG + BM25 (Hybrid)	<ul style="list-style-type: none"> <li>• Semantic + keyword match</li> <li>• Handles diverse queries</li> <li>• No preprocessing</li> <li>• Easy to update index</li> </ul>	<ul style="list-style-type: none"> <li>• Chunking artifacts</li> <li>• Flat hierarchy</li> <li>• May miss context</li> <li>• Reranker adds latency</li> </ul>	Mixed query types, general retrieval
3. RAPTOR (Hierarchical)	<ul style="list-style-type: none"> <li>• Multi-hop reasoning</li> <li>• Handles abstraction</li> <li>• Searches all levels</li> <li>• Good for synthesis</li> </ul>	<ul style="list-style-type: none"> <li>• High preprocessing cost</li> <li>• Summary drift risk</li> <li>• Complex to maintain</li> <li>• Slow index updates</li> </ul>	Conceptual queries, cross-doc synthesis
4. ColPali (Visual Embed)	<ul style="list-style-type: none"> <li>• Preserves layout/tables</li> <li>• No OCR errors</li> <li>• Works on scanned docs</li> <li>• Captures visual context</li> </ul>	<ul style="list-style-type: none"> <li>• Higher compute cost</li> <li>• Larger index size</li> <li>• Less mature tooling</li> <li>• Query needs VLM</li> </ul>	Complex layouts, forms, invoices, scanned PDFs

# Code Examples

## ChromaDB — Simple Vector Store

```
import chromadb

client = chromadb.Client() # in-memory

collection = client.create_collection("docs")
collection.add(
    documents=["doc1", "doc2"],
    ids=["1", "2"]
)

results = collection.query(
    query_texts=["search term"],
    n_results=2
)
```

## txtai — RAG with OpenRouter

```
import os
from txtai import Embeddings, LLM, RAG

os.environ["OPENROUTER_API_KEY"] = "your-key-here"

# LLM via OpenRouter (use litellm/ prefix)
llm = LLM("litellm/openrouter/meta-llama/llama-3.3-70b-instruct")

# Simple generation
response = llm("What is RAG?")
print(response)
```

## RAPTOR — Hierarchical Retrieval with OpenRouter

```
import os
os.environ["OPENROUTER_API_KEY"] = "your-key"

from raptor import (
    BaseSummarizationModel, BaseQAModel, BaseEmbeddingModel,
    RetrievalAugmentationConfig, RetrievalAugmentation
)
from sentence_transformers import SentenceTransformer
from litellm import completion

# Embedding (local, free)
class LocalEmbedding(BaseEmbeddingModel):
    def __init__(self):
        self.model = SentenceTransformer("all-MiniLM-L6-v2")

    def create_embedding(self, text):
        return self.model.encode(text)

# Summarizer via OpenRouter
class OpenRouterSummarizer(BaseSummarizationModel):
    def summarize(self, context, max_tokens=150):
        resp = completion(
            model="openrouter/meta-llama/llama-3.3-70b-instruct",
            messages=[{"role": "user", "content": f"Summarize concisely:\n\n{context}"},],
            max_tokens=max_tokens
        )
        return resp.choices[0].message.content

# QA via OpenRouter
class OpenRouterQA(BaseQAModel):
    def answer_question(self, context, question):
        resp = completion(
            model="openrouter/meta-llama/llama-3.3-70b-instruct",
            messages=[{
                "role": "user",
                "content": f"Answer based on context only.\n\nContext: {context}\n\nQuestion: {question}"
            }]
        )
        return resp.choices[0].message.content
```

```
        }
    )
    return resp.choices[0].message.content

# Setup
config = RetrievalAugmentationConfig(
    summarization_model=OpenRouterSummarizer(),
    qa_model=OpenRouterQA(),
    embedding_model=LocalEmbedding()
)
RA = RetrievalAugmentation(config=config)

# Index document
text = open("your_doc.txt").read()
RA.add_documents(text)

# Query
answer = RA.answer_question("What are the key points?")
print(answer)
```